

May 7, 2026

Hiring Team

Anthropic

Re: AI Policy / Trust & Safety / Risk & Assurance — Anthropic

Dear Hiring Team,

I am writing to express interest in roles at Anthropic that sit at the intersection of AI policy, model risk, and assurance. I have spent 15+ years inside the regulated financial system — at Citi, JPMorgan, and as a Bank Examiner with the Florida Office of Financial Regulation, alongside federal banking regulators — and the last five of those years applying that lens to AI itself. I now run a small consultancy, Audit the Algorithm, and am completing a DBA at FIU on epistemic drift and sycophancy risk in LLM-assisted judgment.

The reason I want to work at Anthropic specifically: your published work on Constitutional AI, responsible scaling, and model evaluation matches the operating discipline I have argued banks and regulators need in order to deploy frontier models safely. I have built that discipline from the inside out — and I would like to build it for the lab whose research most directly informs how the rest of the industry will eventually have to govern these systems.

Relevant track record:

- **Model risk and assurance at scale.** Led CCAR forecast validation and qualitative model challenge for a \$2.6T balance sheet under SR 11-7 and SR 15-18. Contributed to the Fed's "not not-credible" rating on JPMorgan's resolution plan.
- **AI governance in regulated environments.** Designed and shipped a RAG-based Workpaper Quality Assistant inside Citi Internal Audit that cut cycle time 35%, with NIST AI RMF, SR 11-7, XAI, and human-in-the-loop controls built directly into the system.
- **Regulator-facing communication.** Years of writing examiner-defensible workpapers, presenting to audit committees, and translating model behavior into language a regulator will accept. That skill transfers directly to policy, safety case writing, and external assurance work.
- **Practitioner depth on the technology.** I work daily in Python, LangChain/LangGraph, Pinecone/ChromaDB, the Anthropic Claude API, Claude Code, and MCP — both for client work and as part of doctoral research on bias and sycophancy in production-style LLM workflows.

I am confident I can contribute to Trust & Safety, Policy, Model Risk, Assurance, or Responsible Scaling functions, and I am happy to tailor a longer letter to any specific posting. New York or remote both work; I am based in Newark, NJ.

Thank you for your time and your work.

Sincerely,

Yasir A. Malik

YasirAMalik@gmail.com | +1 (786) 704-8536 | [linkedin.com/in/yasiramalik](https://www.linkedin.com/in/yasiramalik) | github.com/AuditingAI