

May 13, 2026

Hiring Team — Internal Audit, Technology & Business Enablement and AI

Citi

388 Greenwich Street, New York, NY

Re: Chief Auditor, Artificial Intelligence — Managing Director (Job Req 26962614)

Dear Hiring Team,

I am applying for the Chief Auditor, Artificial Intelligence role as an internal candidate. The dual mandate this position describes — to audit AI and to build with AI — is the work I have been doing inside Internal Audit at Citi for the past five years, at the scale of a single pilot rather than a global function. I would like to scale it.

The role requires 15+ years of diversified experience across audit, technology, and data-focused mandates. My background covers all three: I have spent the last 15+ years between Citi, JPMorgan, and the U.S. bank regulators (OCC, Federal Reserve, FDIC), and the last five of those translating those audit and risk fundamentals to AI specifically.

Direct evidence against the JD:

- **Independent assurance over AI/GenAI adoption.** Inside Citi IA, I have designed and shipped a RAG-based Workpaper Quality Assistant covering the full AI development lifecycle — use-case intake, model risk tiering, control library, and post-deployment monitoring — with explicit treatment of bias, fairness, hallucination, data privacy, and adversarial-input risks. Audit cycle time fell ~35%. The assistant is itself auditable: NIST AI RMF, SR 11-7, XAI (SHAP/LIME), prompt-injection guardrails, and human-in-the-loop controls are baked into the workflow.
- **Comprehensive coverage of the JD's risk domains.** My prior work spans governance (consent-order audit execution at Citi), legal and regulatory (FR 2900/TIC automation, RRP submissions), model risk (CCAR forecast validation and qualitative model challenge over a \$2.6T balance sheet at JPMorgan, under SR 15-18 and SR 11-7), technology, and third-party — the exact list this role calls for unified treatment of.
- **High-impact, regulator-facing audit delivery.** I led consent-order audit execution at Citi, coordinated cross-functional regulatory deliverables that contributed to the Federal Reserve's "not not-credible" rating on JPMorgan's resolution plan, and have engaged directly with examiners at the OCC, Fed, FDIC, and FL OFR (as a former Bank Examiner myself). I am comfortable in front of Audit Committees and regulators on AI-specific topics.
- **AI-specific frameworks and regulations.** My doctoral research and consulting practice track NIST AI RMF and the GenAI Profile, SR 11-7, ISO/IEC 42001 and 23894, the EU AI Act, and the Colorado AI Act in parallel. My dissertation, *Epistemic Drift & Sycophancy Risk in LLM-Assisted Audit Judgment* (FIU DBA, expected 2028, GPA 3.87), proposes measurement and control mechanisms for one of the unique AI risks this JD identifies.

- **Closing the credential gap.** The JD highlights CISA, CIA, and CRISC as beneficial. I am sitting CIA Part 1 in Q3 2026 and CISA shortly thereafter; the IAPP AIGP (AI Governance Professional) is on the same track.
- **Building and leading elite teams.** I have mentored 8+ auditors on responsible AI adoption — prompt engineering, model risk tiering, and ethical decision-making in human–AI workflows — and was recognized with Citi's *Delivers with Pride* award for AI-powered audit quality assurance innovation. I would be ready on day one to lead both the AI Assurance team and the AI Assurance Tooling & Execution team this role describes.

What I bring that is harder to find than the credentials list suggests: a Bank Examiner's instinct for how a regulator will read a finding, a practitioner's working knowledge of LangChain, vector databases, MLOps awareness, and the Anthropic Claude API, and a researcher's framing for problems the rest of the industry has not yet learned to name. That combination is the one this role needs.

I would welcome the opportunity to discuss how I can build this function for Citi.

Sincerely,

Yasir A. Malik

YasirAMalik@gmail.com | +1 (786) 704-8536 | [linkedin.com/in/yasiramalik](https://www.linkedin.com/in/yasiramalik) | github.com/AuditingAI