

August 3, 2026

Hiring Team — Cloud Responsible AI

Google

New York, NY

Re: Program Manager, Cloud Responsible AI

Dear Hiring Team,

I am applying for the Program Manager role on the Cloud Responsible AI team. The responsibilities in the posting — refining Responsible AI governance frameworks, prioritizing higher-risk product offerings, running pre-launch safety review, and operationalizing adherence to the EU AI Act and ISO 42001 — describe work I have been doing for the past five years, in an industry where getting it wrong is a federal enforcement matter rather than a bad quarter.

Governance frameworks, authored rather than administered. At Citi I wrote the AI governance framework for the internal governance team: use-case inventory, risk tiering so the highest-risk offerings got the most scrutiny, explainability requirements, and human-in-the-loop control design, referenced to NIST AI RMF and SR 11-7. Risk tiering is the same instinct the posting names — deciding where the standard has to be strictest before anyone asks.

Pre-launch review under adversarial conditions. I led consent-order remediation, preparing the issue-closure evidence submitted to quality assurance and external regulators. An institution under a federal order has every incentive to declare itself fixed; my job was determining whether the evidence supported the claim, and writing findings that held when the subject was powerful and motivated to disagree. That is what pre-launch safety analysis demands: a defensible no, delivered to people who wanted a yes.

I have also shipped the thing being governed. I built a retrieval-augmented document assistant that cut review cycle time about 35% — requirements through delivery and adoption. So when a framework meets an engineering team, I know which requirements are real constraints and which are theater. Governance written by someone who has never shipped tends to get routed around.

The human-factors half. My doctoral research at FIU examines bias in professional judgment: I completed an IRB-approved empirical study in July 2026, and my dissertation in development extends it to automation bias — when professionals stop exercising judgment because a system sounded confident. Enterprise customers adopting Cloud AI will fail less on model quality than on how their people use the output. That is a research question I am actively working on, not a talking point.

Where I am honest about the gap. My background is regulated financial services, not consumer trust and safety, and I have not worked on a Cloud or API product in a B2B context. What I bring instead is five years of enterprise AI risk governance, twenty years of program delivery against regulators, and a bank examiner's instinct for how an outside reviewer will read a control. The vocabulary of a new domain is learnable; the judgment about when a system is safe enough to launch is the part that takes years.

I would welcome the conversation.

Sincerely,

Yasir A. Malik

Newark, NJ · YasirAMalik@gmail.com · +1 (786) 704-8536