

September 19, 2026

Hiring Team — Core AI Foundations

Google

New York, NY (preferred working location) · San Jose, CA · Redwood City, CA

Re: Vice President, Product Management, Core AI Foundations

Dear Hiring Team,

I am writing about the Core AI Foundations role, and I will start where your screen will start, because you will reach it in ten seconds and I would rather you hear it from me.

I have not held a product-manager title, and my degree is in Banking and Finance, not Computer Science. Your stated minimum is the full scope of product management at a depth I do not have on paper, and if that is the binding constraint, this is a short letter and I understand it completely.

What I do have is the half of a foundation-model organisation that decides whether a model can leave the building. At Citi I authored the AI governance framework — NIST AI RMF and SR 11-7 — and defined the AI tooling roadmap it governed: use-case intake, risk tiering by criticality, a control library, explainability and human-in-the-loop requirements, and post-deployment monitoring KPIs. That is a product roadmap by another name: what gets built, in what order, and what has to be true before it ships.

Then I built the thing the framework governed, which is the part that is unusual. I elicited requirements from the auditors who needed a workpaper assistant, wrote it in Python and LangChain with vector search, SHAP/LIME explainability and prompt-injection guardrails, integrated it with Alteryx, iterated on their feedback, and put it into production — **~35%** faster review cycles. Governance people who cannot build tend to write policy nobody can implement; builders without governance ship things that cannot be defended. I have done both jobs.

The regulated-market problem is the one I know from the inside. I was a bank examiner at the Florida Office of Financial Regulation before I was ever audited by one. Later I led consent-order remediation where issue-closure evidence had to be accepted by federal regulators, and ran Resolution & Recovery Planning across **50+ stakeholders** in five functions — none reporting to me — into Federal Reserve and FDIC submissions against statutory deadlines. As foundation models move into finance, health and public infrastructure, the binding constraint stops being capability and becomes defensibility. That is the question I have spent my career on.

My doctoral research is directly on your failure mode. My DBA at FIU measures automation bias in AI-assisted judgment: how expert judgment degrades when a model is confidently wrong. The completed study (IRB-25-0462) was a 55-item instrument designed and fielded under IRB approval; the dissertation in development extends it to AI assistance. Sycophancy and over-reliance are usually discussed as evaluation problems. They are also human-factors problems, and they are measurable.

I am not going to claim I could walk into a VP product seat over Google's core model stack. If you need someone who has run a large product organisation through several model generations, that is not me. If part of what this organisation needs is someone who can hold the trust, evaluation and regulated-deployment surface — and write code and read a regulator — I would welcome a conversation about where that fits, including below the level posted.

Sincerely,

Yasir A. Malik

YasirAMalik@gmail.com · +1 (786) 704-8536 · Newark, NJ

linkedin.com/in/yasiramalik · github.com/MalikAI-786