



YASIR A. MALIK

Curriculum Vitae

Newark, NJ • YasirAMalik@gmail.com • +1 (786) 704-8536 • linkedin.com/in/yasiramalik • github.com/MalikAI-786 • auditingai.github.io/research.html

EDUCATION

Doctor of Business Administration (DBA), Florida International University — expected 2028

Chapman Graduate School of Business, Miami, FL • GPA 3.81 • in progress

- Qualifying examination passed July 2026. Qualifying research: *Mitigating Anchoring Bias in Long-Term Auditor Engagements* — IRB-approved instrument (IRB-25-0462) built and fielded.
- Dissertation in development: automation bias and judgment drift when auditors work with AI assistants. Not yet under ethics review; no participants approached.

MBA, Financial Mathematics, Florida International University — 2011

GPA 3.8 • preceded by the Pre-MBA program in International Banking, FIU, 2007

B.Sc., Banking & Finance, London School of Economics and Political Science — 2005

Columbia Engineering FinTech Boot Camp, Columbia University — 2021

DOCTORAL RESEARCH PROGRAMME — *Auditing the Auditor*

What happens to professional judgment when the machine answers first — and agrees with you. Every framework now governing AI (SR 11-7, the NIST AI Risk Management Framework, the EU AI Act) assumes a competent human reviewer sits above the model; none measures whether that reviewer's judgment survives contact with it. The programme separates three links that are usually conflated, of which only the first is a cognitive bias:

Link	Claim	Construct	Question · method
L1 · Automated anchoring	System output functions as an anchor — automated, continuous, arriving before the reviewer has formed a view	Automation bias — a human cognitive bias	How much · survey
L2 · Sycophantic confirmation	The model agrees with a stated position rather than challenging it	Sycophancy — model behaviour, not a bias	How · interviews
L3 · Recursive epistemic drift	Successive systems reprocess earlier machine-influenced work and converge on each other rather than on evidence	Model collapse — a property of a system of models	How, over time · longitudinal

Status, stated plainly

Element	State
Qualifying examination	Passed, July 2026
Model — eleven constructs, fifty-five items, sixteen hypotheses	Built
Instrument — IRB-approved (IRB-25-0462), fielded	Built and fielded
Hypotheses tested	None — the sample did not support it (below)
Peer-reviewed publications	None to date
The AI extension (L1–L3)	Argued, not tested — no findings exist

The finding that came out of the failure

The instrument worked; the population did not exist. Applying the study's eligibility criteria to a commercial research panel in July 2026: **334,976** panel members screened, **~20** matched the eligibility criteria (a prevalence near six per hundred thousand), **23** raw responses recorded, **4** survived screening. A survey at that prevalence needed a sampling frame of roughly 9.6 million; the panel held about three and a half percent of it. This is a structural property of studying a narrow specialist population, measurable in advance, which the profession's methods literature treats as a limitation to apologise for rather than a parameter to design around. It is now a research question in its own right: under what conditions is a cross-national study of a rare professional population feasible at all?

Open gaps, including the uncomfortable ones

- **Measurement.** The outcome construct is self-reported — it asks auditors to report how far a reference point drove their judgment, which is precisely what anchoring prevents them from noticing. Needs an honest rename to *perceived judgment discipline*, or a behavioural measure: a planted anchor and an observable adjustment.
- **Theory.** Automation bias and algorithm aversion both have support; L1 assumes over-trust and has no account of when the reverse occurs. The moderators need specifying — task type, expertise, stakes, and whether output arrives before or after the human forms a view.
- **Access.** L3 has no design, deliberately: it needs multi-year engagement files or repeated access to the same auditors. Writing a protocol for a study that cannot be run would repeat the error already made.

Current work

- **Feasibility manuscript** — **drafted, not submitted.** The prevalence finding written as a methodological contribution for cross-national research on hard-to-reach professional populations.
- **Feasibility calculator** — **working.** Takes panel size, prevalence, eligibility and completion rates; returns reachable sample and cost per usable response; validated by reproducing the study that produced it. Its useful output is the threshold where a survey design stops being viable and an interview design becomes the honest choice.
- **The qualitative arm (L2)** — **blocked on ethics review.** Phenomenological design: protocol, sampling plan, coding scheme, trustworthiness criteria and an append-only audit trail written. Falsification conditions committed and timestamped *before* any data exists. Zero participants.
- **A research pipeline built against its own subject.** Model-agnostic runbooks: every claim carries a resolving source, a null search is reported as null, and an adversarial pass runs against the programme's own conclusions — because a researcher who lets a model confirm his topic is running his own dissertation's failure mode on himself.

TEACHING, TUTORING & MENTORING

[TO CONFIRM: Graduate Teaching Assistant — institution, department, course title(s) and number(s), term(s), supervising faculty member, and duties (sections led, grading, office hours).]

[TO CONFIRM: SAAC Tutor — spell out the acronym, institution, subjects tutored, dates, and whether one-to-one or group.]

Guest lecturer, Florida International University — ongoing

[TO CONFIRM: course(s), host faculty, topic(s) and date(s) of guest lectures.]

Class visit, Rutgers University — 15 Sep 2026

[TO CONFIRM: campus, course, host, and whether this is a talk or an observation.]

Mentor to 8+ internal auditors, Citi Internal Audit — 2021 – 2026

- Responsible AI adoption, prompt engineering, and ethical decision-making in human–AI workflows; recognised with Citi's *Delivers with Pride* award.

Teaching interests

Auditing and assurance • internal audit and controls • enterprise and operational risk • bank regulation and supervision • AI governance and AI ethics • business analytics and data governance • fintech • judgment and decision-making in professional practice

PROFESSIONAL EXPERIENCE

Vice President, Audit Manager — Citi — Jul 2021 – Apr 2026

Internal Audit, Cross-Enterprise Program & Change Management, New York

- Led risk-based audits across 15+ business units; consent-order issue-closure evidence accepted by external regulators; built independent assurance over enterprise AI adoption; designed and shipped a RAG workpaper assistant (Python, LangChain) that cut review cycle time ~35%; executive and Board Audit Committee reporting.

Risk Control Manager, Treasury & CIO — JPMorgan Chase — Mar 2019 – Jun 2021

First-line capital controls and Resolution & Recovery Planning, Jersey City

- Owned the first-line control framework for capital activities; CCAR forecast validation over a \$2.6T balance sheet under SR 15-18; Resolution & Recovery Planning across 50+ stakeholders into Federal Reserve and FDIC submissions; automated reconciliations, ~40% less manual review.

Capital Controller, Basel Measurement & Analytics — JPMorgan Chase — Sep 2017 – Feb 2019

Brooklyn; preceded by Program Manager, CIB Resolution & Recovery Planning, 2015–2017

- Basel III RWA and capital adequacy reporting for a \$50B book; \$180M in capital optimization identified.

Assistant Vice President, Global Legal Entity Management — Citi — 2012 – 2015

Tampa

- Consolidated 500+ legal-entity data sources into governed master data; automated FR 2900 and TIC regulatory filings at 99.8% accuracy.

Bank Examiner, Bureau of Bank Regulation — Florida Office of Financial Regulation — Apr 2011 – Mar 2012

West Palm Beach

- CAMELS safety-and-soundness examinations of state and national banks alongside federal banking regulators; workpapers supporting formal enforcement actions.

Senior Business Analyst, Retail Credit Risk — Royal Bank of Scotland — 2008 – 2009

Dubai

- Credit-risk MIS over a \$500M+ retail portfolio in SAS, SQL, and Excel/VBA.

INDUSTRY PRACTICE

Founder, Audit the Algorithm — 2024 – present

- AI governance advisory for regulated financial services (auditingai.github.io); bias and drift testing, regulator-defensible control frameworks, human–AI workflow design. Authored an AI governance framework referencing NIST AI RMF and SR 11-7.

RESEARCH METHODS & SKILLS

Methods: survey instrument development and validation • experimental design • IRB protocol authorship • feasibility and prevalence analysis for hard-to-reach populations • phenomenological interview design and coding • pre-registered falsification conditions • adversarial review of one's own conclusions

Statistical and analytical software: SPSS • SAS • Python (scikit-learn, pandas) • SQL • Excel/VBA

AI and NLP: LangChain / LangGraph • retrieval-augmented generation • vector databases • SHAP / LIME explainability • evaluation and red-teaming

Visualisation and workflow: Tableau • Power BI • Alteryx • GitHub Actions

CERTIFICATIONS & TRAINING

FDIC Bank Examiner I • GCP Social/Behavioral Human Subjects Research • Registered Scrum Master. *In progress:* Certified Internal Auditor (Part 1) • IAPP Artificial Intelligence Governance Professional (AIGP).

RECOGNITION

Citi *Delivers with Pride* award for AI-powered audit quality and mentoring.

REFERENCES

Available on request.