



YASIR A. MALIK

Founder, Audit the Algorithm | AI Governance Advisory for Regulated Financial Services

New York metro • YasirAMalik@gmail.com • +1 (786) 704-8536 • auditingai.github.io • linkedin.com/in/yasiramalik

THE PROBLEM I SOLVE

AI-assisted review can pass every model validation and still degrade the judgment it was meant to support. The model checks out; the people using it quietly stop checking. Most firms have governance for the first failure and nothing for the second — and since April 2026 the regulator's own model risk guidance, **SR 26-2**, puts generative and agentic AI **expressly outside its scope**. The control expectations for those systems are not yet written. I help firms write theirs before an examiner asks, from the perspective of someone who has been the examiner, the first line, and the auditor.

WHAT I DO

AI governance, built to be examined	Frameworks a regulator can test: use-case intake, risk tiering, control library, explainability and human-in-the-loop requirements, post-deployment monitoring. Mapped to NIST AI RMF, the EU AI Act, ISO/IEC 42001 and SR 26-2.
Judgment risk in AI-assisted review	Separates three failures that usually get blurred into one: people over-trusting the tool, the model agreeing with whoever is asking, and systems converging on each other instead of on evidence. Each gets its own control.
Controls you can measure	Turns commonly claimed audit-quality controls into something a firm can actually measure rather than assume — a 55-item instrument covering eight intervention categories, designed and fielded under institutional ethics approval.
Go / no-go before the spend	A calculator that prices a study's or programme's weakest input before anyone signs a vendor contract. Proven on my own research: a panel of 334,976 held about six eligible people per hundred thousand — a failed survey, caught before it became a failed budget.
Build, not only advise	Retrieval-augmented tools in Python and LangChain with explainability and prompt-injection guardrails, built to be auditable themselves — and a research process that does not quietly agree with whoever is paying for it.

WHERE IT WAS PROVEN

- **Citi, Internal Audit, New York** — authored the enterprise AI governance framework and built and shipped the retrieval assistant it governed: ~35% faster audit review. Risk-based audits across **15+ business units**; consent-order remediation accepted by external regulators; Board Audit Committee reporting.
- **JPMorgan Chase** — CCAR forecast validation over a **\$2.6T** balance sheet; Basel III analysis that surfaced **\$180M** in capital optimisation; Resolution & Recovery Planning across **50+ stakeholders** into Federal Reserve and FDIC submissions.
- **Citi, Global Legal Entity Management** — **500+** data sources consolidated into governed master data behind FR 2900 and TIC filings at **99.8%** accuracy.
- **Florida Office of Financial Regulation** — bank examiner in the aftermath of the 2008 crisis: CAMELS examinations alongside federal banking regulators, workpapers supporting formal enforcement actions.
- **Royal Bank of Scotland, Dubai** — retail credit-risk MIS for a **\$500M+** portfolio through the 2008 crisis, as the bank itself went into government rescue.

HOW I WORK

Evidence before assertion: nothing claimed without a primary source, open questions flagged rather than smoothed over, and the record kept of what did not work as well as what did. The public working record is at auditingai.github.io — evidence of how the work is actually done, not a portfolio of finished things.

EDUCATION & CREDENTIALS

Doctor of Business Administration, Florida International University — in progress; qualifying examination passed July 2026. Research: how AI-assisted review degrades professional judgment. • **MBA, Financial Mathematics, FIU** — 2011 • **B.Sc., Banking & Finance, London School of Economics** — 2005

Certifications: FDIC Bank Examiner I • Registered Scrum Master; *in progress* CIA Part 1, IAPP AIGP