



YASIR A. MALIK

AI Governance & Model Risk | Responsible AI at Enterprise Scale | Ships the Systems as Well as the Policy | Doctoral Researcher in Human-AI Judgment
Newark, NJ • Preferred working location: New York, NY • YasirAMalik@gmail.com • +1 (786) 704-8536 • linkedin.com/in/yasiramalik • github.com/MalikAI-786

SUMMARY

Works on the half of an AI platform that decides whether a model can be deployed into a regulated market and defended afterwards. Authored an **AI governance framework** at Citi referencing **NIST AI RMF** and **SR 11-7** — use-case intake, risk tiering, control library, explainability and human-in-the-loop requirements, post-deployment monitoring KPIs — and defined the AI tooling roadmap it governed. Then built what it governed: requirements elicited from the teams who needed it, a retrieval-augmented assistant written in **Python** and **LangChain** with vector search, SHAP/LIME explainability and prompt-injection guardrails, iterated to production — **~35%** faster review cycles. The career before AI was capital and regulatory: **CCAR** validation over a **\$2.6T** balance sheet, Basel III RWA analysis surfacing **\$180M** in capital optimisation, and Resolution & Recovery Planning across **50+ stakeholders** into Federal Reserve and FDIC submissions. A former bank examiner, so the regulator's read on a launch decision is a learned instinct. **DBA candidate at FIU** researching automation bias in AI-assisted judgment — the over-reliance problem, measured rather than asserted.

WHAT A FOUNDATION-MODEL PRODUCT ORGANISATION NEEDS, AGAINST THE RECORD

What the work requires	Where it has been done
Product roadmap ownership and prioritisation	Defined the AI tooling roadmap for Citi Internal Audit: use-case intake, risk tiering by criticality, control library, and post-deployment monitoring KPIs — what gets built, in what order, and what has to be true before it ships.
Requirements to shipped product, with users in the loop	RAG Workpaper Quality Assistant: requirements elicited from the auditors who needed it, written in Python and LangChain with vector search, SHAP/LIME explainability and prompt-injection guardrails, integrated with Alteryx, iterated on feedback, and put into production. ~35% faster review cycle.
Launch gating, evaluation and model risk	Independent assurance over enterprise AI adoption — governance, model risk, bias and fairness, hallucination, adversarial robustness, third-party AI; SR 11-7; SHAP/LIME; evaluation and red-teaming scripts. Consent-order remediation with evidence accepted by external regulators: pre-launch review under the least forgiving reviewer there is.
Deploying into regulated markets — EU AI Act and after	NIST AI RMF, EU AI Act, ISO/IEC 42001. Former bank examiner (Florida Office of Financial Regulation), CAMELS examinations alongside federal banking regulators — the supervisory view of a product decision, from inside the supervisor.
The human side of model quality	DBA research at FIU on automation bias in AI-assisted judgment : how expert judgment degrades when a model is confidently wrong. Qualifying study completed Jul 2026 (IRB-25-0462), a 55-item instrument designed and fielded under IRB approval. The sycophancy and over-reliance problem, measured.

TECHNICAL

Build: Python • LangChain / LangGraph • retrieval-augmented generation and agentic RAG • vector databases (Pinecone, ChromaDB) • Hugging Face • scikit-learn • SHAP / LIME • evaluation and red-teaming scripts • prompt engineering and injection guardrails • SQL • GitHub Actions

Govern: NIST AI RMF • EU AI Act • ISO/IEC 42001 • SR 11-7 model risk • Global Internal Audit Standards • COSO • Three Lines of Defense

Analyse: SPSS (experiment design and statistical testing) • SAS • Alteryx • Tableau • Power BI • Excel/VBA. **Delivery:** Registered Scrum Master.

PROFESSIONAL EXPERIENCE

Vice President, Audit Manager — Citi — Jul 2021 – Apr 2026

Internal Audit, Cross-Enterprise Program & Change Management | New York, NY

- **Authored the AI governance framework** referencing NIST AI RMF and SR 11-7, and defined the AI tooling roadmap it governed — use-case intake, risk tiering, control library, explainability and human-in-the-loop requirements, post-deployment monitoring KPIs.
- **Designed and shipped the RAG Workpaper Quality Assistant** — Python, LangChain, vector search over unstructured workpapers, SHAP/LIME explainability, prompt-injection guardrails, Alteryx integration. Specified with its users, iterated on their feedback, put into production: **~35%** less review cycle time, and continuous coverage where sampling had been the only option.
- Built independent assurance over enterprise AI adoption across governance, model risk, bias and fairness, hallucination, adversarial robustness, and third-party AI; led risk-based audits across **15+ business units**; drove consent-order remediation with evidence accepted by external regulators. Executive and Board Audit Committee reporting; coached **8+ auditors**.

Risk Control Manager, Treasury & CIO — JPMorgan Chase — Mar 2019 – Jun 2021

First-Line Capital Controls & Resolution and Recovery Planning | Jersey City, NJ

- Owned the first-line control framework for Treasury and CIO capital activities; **automated capital and liquidity reconciliations**, cutting manual review **~40%** and removing a recurring source of reporting error. Led **CCAR** forecast validation and quantitative model challenge over a **\$2.6T** balance sheet under SR 15-18.
- Ran **Resolution & Recovery Planning** across **50+ stakeholders** in five functions — none reporting to him — into Federal Reserve and FDIC submissions against statutory deadlines.

Capital Controller, Basel Measurement & Analytics — JPMorgan Chase — Sep 2017 – Feb 2019

Brooklyn, NY • preceded by Program Manager, CIB Resolution & Recovery Planning (2015–2017)

- Owned Basel III RWA and capital adequacy reporting across a **\$50B** book of equities, fixed income and OTC derivatives; analysis of RWA treatment surfaced **\$180M** in capital optimisation for CFO decision support.

Assistant Vice President, Global Legal Entity Management — Citi — 2012 – 2015

Tampa, FL

- Consolidated **500+ legal-entity data sources** into governed master data with ownership, lineage and data-quality controls — enabling automated **FR 2900** and **TIC** filings at **99.8% accuracy** under regulatory examination.

Bank Examiner, Bureau of Bank Regulation — Florida Office of Financial Regulation — Apr 2011 – Mar 2012

West Palm Beach, FL • earlier: Senior Business Analyst, Retail Credit Risk — Royal Bank of Scotland, Dubai (2008–2009)

- Conducted CAMELS safety-and-soundness examinations of state and national banks alongside federal banking regulators; authored workpapers supporting formal enforcement actions.

EDUCATION, RESEARCH & ADVISORY

Doctor of Business Administration (DBA), Florida International University — in progress. Qualifying examination passed Jul 2026 (IRB-25-0462): a 55-item instrument on anchoring in audit judgment, designed and fielded under institutional ethics approval. Dissertation in development: *Anchoring Bias in LLM-Assisted Audit Judgment*.

MBA, Financial Mathematics, Florida International University — 2011 • **B.Sc., Banking & Finance, London School of Economics** — 2005 • **Columbia Engineering FinTech Boot Camp** — 2021

Founder, Audit the Algorithm (auditingai.github.io, 2024–): AI governance advisory for regulated financial services. **Certifications:** FDIC Bank Examiner I • Registered Scrum Master • GCP Social/Behavioral Human Research. *In progress:* CIA Part 1, IAPP AIGP.