

Yasir A. Malik

Newark, New Jersey · _____ · _____

BMO U.S. — Head, Responsible AI & Governance

Dear Hiring Team,

I am writing regarding the Head, Responsible AI and Governance role for BMO's U.S. business.

I am applying to build the function rather than to review it. Fifteen years in audit and risk at JPMorgan Chase and Citigroup, preceded by examinations work at the OCC and a U.S. state regulator, has left me with a specific skill: I know what a governance function has to produce before an examiner asks for it, and I know the difference between a policy that exists and a policy that holds. Those are not the same thing, and the gap between them is where most AI governance programs are currently sitting.

The build I would bring is concrete. At Citigroup I authored an AI governance framework referencing the NIST AI Risk Management Framework and SR 11-7 — risk tiering, explainability requirements, and human-in-the-loop controls scaled to the tier rather than applied uniformly. I also built the technology it governs: a retrieval-augmented AI review assistant deployed into the audit function, which cut review cycle time by roughly 35%. Having shipped one of these systems changes what you write into a standard. You stop writing controls for the model and start writing them for the workflow the model sits inside, which is where the failures actually occur.

On the supervisory side, I have run the kind of work a head of this function has to be able to defend. I led consent-order remediation at Citigroup where the evidence package was accepted by federal regulators. At JPMorgan Chase I owned Basel III risk-weighted asset and CCAR reporting — FR Y-14Q, FFIEC 101, Pillar 3 — and led Resolution and Recovery Planning across more than fifty stakeholders for Federal Reserve and FDIC submissions. Thematic audits across more than fifteen Citigroup business units, with the reporting structured for the Board Audit Committee, is the same mechanic responsible AI reporting needs at the enterprise level.

I would note one thing about BMO's shape specifically. A U.S. platform inside a Canadian parent means the AI governance framework has to satisfy U.S. supervisory expectations while mapping cleanly onto the home supervisor's model risk and third-party standards. My regulatory experience is U.S. — OCC and a state regulator — so I would not claim familiarity with the Canadian regime itself. What I would claim is that building a control framework that answers to two supervisors without maintaining two frameworks is a design problem I have worked on before, and it is one that gets much more expensive if it is deferred.

My doctoral research at Florida International University, where I passed the qualifying examination in July 2026, examines automation bias in professional judgment — what happens to an expert's conclusion when a confident machine output arrives early in their process. I raise it because it is the

exposure a model validation program does not catch: the model can perform exactly to specification and the decision still degrades. For a bank deploying AI into credit, compliance, and customer-facing judgment, that belongs in the control design from the start.

I would welcome a conversation about how the function is being scoped.

Sincerely, Yasir A. Malik _____ · _____