Yasir A. Malik

Newark, New Jersey · ____________________ · ____________________

MetLife — VP, AI Risk & Governance

Dear Hiring Team,

I am writing regarding the VP, AI Risk and Governance role in Bridgewater. Bridgewater is a routine commute from where I live in Newark, and I am available onsite on whatever schedule the team runs.

My model governance experience is from banking. I want to be straightforward about what that means for an insurance role: I have not worked under NAIC model governance expectations, I am not an actuary, and I have not been examined by a state insurance department. What I have done is spend fifteen years in audit and risk at JPMorgan Chase and Citigroup, and before that examinations work at the OCC and a U.S. state regulator, in an environment where every model that touched a reported number had to be inventoried, tiered, documented, independently challenged, and defended to an outside examiner who was not inclined to accept the first answer. That discipline is the same discipline whether the model is pricing credit risk or pricing mortality — what changes is the regulator, the vocabulary, and which committee owns the sign-off.

The specifics that transfer most directly. At Citigroup I authored an AI governance framework referencing the NIST AI Risk Management Framework and SR 11-7 — risk tiering, explainability requirements, and human-in-the-loop controls calibrated to tier. NIST AI RMF is regulator-neutral and is the framework I would build from here; the banking guidance I would treat as one source of prior art among several rather than as the anchor. I also ran thematic audits across more than fifteen Citigroup business units and structured the reporting that went to the Board Audit Committee, and I led consent-order remediation where the evidence package was accepted by federal regulators. Enterprise model inventory, lifecycle gates that actually stop something, and evidence that survives an external party pulling on it — that is the mechanic, and it is portable.

On the technology itself: I built and deployed a retrieval-augmented AI review assistant inside Citi's audit function that cut review cycle time by roughly 35%. Having shipped one changes how you write the controls. You stop governing the model in isolation and start governing the workflow it sits inside — which for a carrier means underwriting, claims, and actuarial review, where the professional making the decision is the control.

That is also what my doctoral research is about. At Florida International University, where I passed the qualifying examination in July 2026, I study automation bias in professional judgment — what happens to an expert's conclusion when a confident machine output arrives early in their process. For a business built on actuarial and underwriting judgment, that is the AI exposure a validation program does not catch, because it appears even when the model is performing exactly to specification.

I would welcome a conversation about how the function is structured and what the first year looks like.

Sincerely, Yasir A. Malik ____________________ · ____________________