

Website Content Kit — Yasir A. Malik / Audit the Algorithm

Rule for everything below: every claim is verifiable. Items in [BRACKETS] need your confirmation before publishing. Nothing here credits an AI vendor or implies a partnership that doesn't exist — the strength of the positioning is that *you* did the work.

1 · Positioning — the one-sentence version

Three things almost nobody in AI governance can claim together: I examined banks as a regulator, I ran audits inside two of them, and I built and shipped an AI tool into a live audit function. Now I research where that technology quietly breaks professional judgment.

Shorter, for a hero line:

Bank examiner. Audit executive. AI builder. Now researching how AI reshapes professional judgment — and what governance actually has to catch.

Bio (100 words, for speaker pages and LinkedIn):

Yasir A. Malik is a doctoral researcher and practitioner in audit, risk, and AI governance. Over fifteen years he led internal audit and risk work at Citigroup and JPMorgan Chase, following service as a bank examiner with the OCC and state regulators. At Citi he built and deployed a retrieval-augmented AI assistant into the audit function, cutting review cycle time by roughly 35%, and proposed an AI governance framework grounded in the NIST AI Risk Management Framework and SR 11-7. His doctoral research at Florida International University examines cognitive bias in audit judgment — and what happens to that judgment when AI systems enter the workflow.

2 · The research — what actually exists, stated accurately

Completed and passed (July 2026): DBA qualifying examination, Florida International University.

Mitigating Anchoring Bias in Long-Term Auditor Engagements — an eleven-construct, fifty-five-item measurement model with sixteen hypotheses, grounded in dual-process theory and the anchoring-and-adjustment heuristic. IRB-approved (IRB-25-0462), instrument fielded, complete analytic protocol specified from exploratory factor analysis through structural modeling.

The model in plain language: eight organizational practices that should disrupt anchoring — training and awareness, auditor rotation, analytical tools, structured processes, feedback and reflection, independent review, regulatory guidance, and performance incentives — operating through two mechanisms (the quality of the auditor's judgment, and the rigor of the audit process) toward one outcome: judgments driven by current evidence rather than last year's number.

Faculty assessment, quotable:

"Your reflection on the challenges encountered during the study... demonstrates sound research judgment." — Dr. Juan C. Rey, FIU College of Business, final evaluation of the qualifying examination.

Honest framing of the empirical result (use this — it is more credible than hiding it, and it *is* the finding):

The study reached feasibility, not validation. Recruiting the target population proved structurally impossible: applying the eligibility criteria to a research panel of over 330,000 participants returned roughly twenty eligible people — about six in one hundred thousand. That result reframed the work as a feasibility study and produced its

most useful contribution: evidence that specialist professional populations cannot be reached through standard research infrastructure at any price.

Current research direction (in development — say "in development", not "findings"):

How AI and large language models interact with professional judgment in recurring review work. The chain under investigation: system-generated output becomes the anchor itself — automated, continuous, at scale; the model then tends to confirm the position a user has already stated rather than challenge it; and as successive models reprocess the same work, conclusions converge on each other rather than on evidence.

3 · Services / partnering — what you can credibly offer today

Audit the Algorithm — governance-first AI advisory.

Offer	What it is	Who buys it	Why you
AI governance readiness review	Assess an organization's AI use against NIST AI RMF and, for models, SR 11-7 — inventory, documentation, human-oversight design, evidence trail	Banks, insurers, asset managers, regulated mid-market	You have proposed exactly this framework inside a G-SIB
Human-oversight design for AI-assisted review	Where reliance is calibrated correctly and where it quietly isn't — reviewer workflow, challenge points, documentation standards	Internal audit, risk, compliance functions adopting AI tools	This is literally your research question
Audit-function AI enablement	Scoping and building AI assistance into audit workflow, with controls designed in from the start	CAEs and audit directors under efficiency pressure	You built and shipped one — 35% cycle-time reduction
Regulator-readiness translation	Preparing AI-related documentation and narratives for examination scrutiny	Firms anticipating supervisory questions on AI	Former examiner: you have sat on the other side of the table
Workshops / guest lectures	Cognitive bias in professional judgment; AI over-reliance and automation bias; anchoring in recurring engagements	Professional bodies, firm training programs, universities	Practitioner credibility plus doctoral research

Partnering angles worth pursuing (each is a real, checkable opening — not a claim):

- 1 Professional bodies** — IIA, ISACA, AICPA chapters: content partnership on AI governance for practitioners. They need practitioner-researchers; you are one.
- 2 Universities** — guest lecture, adjunct, or case-material collaboration. Your Rutgers and FIU paths are already open.
- 3 Audit software / GRC vendors** — advisory or research collaboration on responsible-AI features. They need credibility with regulated buyers.

- 4 **Research collaboration** — your dissertation needs a large professional sample; firms and associations need to understand AI risk in their own functions. That is a genuine two-way trade: they get findings about their people, you get the sample.
 - 5 **Regulator and standard-setter engagement** — comment letters and working groups on AI in audit (PCAOB and IAASB are both actively consulting).
-

4 · Proof points (all verifiable — no embellishment)

- Fifteen years in audit and risk at **Citigroup** and **JPMorgan Chase**
 - **Bank examiner** — OCC and Florida Office of Financial Regulation [CONFIRM final agency wording]
 - Built and deployed a **production RAG assistant** in Citi's audit function — **~35% review cycle-time reduction**
 - Proposed an **AI governance framework** referencing **NIST AI RMF** and **SR 11-7** for an internal governance team
 - Led the firm's **Resolution and Recovery Planning** program at JPMorgan, coordinating 50+ stakeholders on Federal Reserve and FDIC submissions
 - Owned **Basel III RWA** calculation and capital-adequacy reporting; identified **\$180M** in capital optimization
 - **DBA qualifying examination passed**, FIU, July 2026 — [CONFIRM expected completion year: official 2024–25 evaluation record indicates 2028 defense; resume states 2027]
 - **MBA**, Florida International University, 2011 · **BSc Banking & Finance**, London School of Economics, 2005
 - **Certifications:** FDIC Bank Examiner I · Registered Scrum Master · CITI human-subjects research (social/behavioral) [CONFIRM current list]
-

5 · How to talk about AI on the site (this matters)

Do say:

- "I build with these tools daily and I research where they fail."
- "My work uses AI for research assistance, analysis, and drafting support — the judgment, synthesis, and interpretation are mine."
- Name tools as tools when relevant (Qualtrics, Jamovi, Python, LLM assistants), the way any researcher names instruments.

Do not say:

- "Validated by / partnered with / powered by" any AI vendor. You have no such relationship, and the claim would be checked.
- Anything implying AI co-authorship of the research. Your faculty explicitly asked students to ensure their scholarly voice is their own; your site should reflect that standard.

The strongest honest line about your AI use:

I use AI systems intensively — which is precisely why I study their failure modes. The risk I research is not hypothetical to me; it is the risk in my own workflow, managed deliberately.

6 · Suggested site structure

- 1 **Home** — the positioning line, three proof points, one call to action
- 2 **Research** — the qualifying study, the model diagram, the current direction, the honest feasibility finding
- 3 **Advisory** — the five offers above, with the regulated-industry framing
- 4 **Writing / Speaking** — publications as they come, workshop topics now
- 5 **About** — the bio, the examiner-to-executive-to-researcher arc
- 6 **Contact** — one form, one email

Assets you already have: the research model figure, the dissertation-extension figure, the study overview document, the full manuscript, and the public project repository (a genuine differentiator — very few consultants can show their work in a timestamped public record).

7 · Before publishing — verification checklist

- ☐ DBA completion year — reconcile 2027 vs 2028 against the program office
- ☐ GPA if stated anywhere — official record shows 3.81, resume states 3.87
- ☐ Exact examiner agency and dates
- ☐ Current certification list
- ☐ Whether "Audit the Algorithm" is a registered entity or a working name — say "advisory practice" if not yet formed
- ☐ Any client work must be described without identifying employers' confidential detail
- ☐ Nothing implies the dissertation is complete, published, or that its extension has findings yet