



Receiving Confirmation

How experienced auditors make sense of AI-generated agreement with a judgment they had already formed

RESEARCH QUESTION

How do experienced auditors experience and make sense of receiving an AI-generated conclusion that confirms a judgment they had already formed?

1 · PROBLEM STATEMENT

An auditor on a recurring engagement has always carried a prior conclusion forward. What is new is that an AI system now produces a conclusion of its own, and sometimes it agrees.

Research on AI and auditor judgment has focused primarily on the extent of reliance, which a survey can measure. The lived experience of receiving machine agreement with a judgment already formed has received much less attention.

Firms and regulators are writing AI-use policy on the assumption that a human reviewer still exercises independent judgment.

2 · PURPOSE STATEMENT

The purpose of this phenomenological study is to understand how experienced auditors experience and make sense of receiving an AI-generated conclusion that confirms a judgment they had already formed.

3 · PHILOSOPHICAL ASSUMPTIONS

SOCIAL CONSTRUCTIVIST · INDUCTIVE ONTOLOGICAL

Multiple realities. Auditors' accounts of the experience are the data, not a proxy for one correct account.

EPISTEMOLOGICAL

Knowledge comes from first-person accounts in semi-structured interviews.

AXIOLOGICAL

The researcher is a career auditor who built an AI review tool: insider access and insider bias, managed through reflexive memos and a confirmation hazard log.

METHODOLOGICAL

Inductive phenomenology across individuals at different firms.

4 · METHODOLOGY: APPROACH

Phenomenology (Creswell & Poth, 2024). The question asks about experience and meaning making across individuals, which is what phenomenology is for.

- Not case study: the phenomenon crosses firms, with no bounded site
- Not grounded theory: the question is about experience, not process

5 · METHODOLOGY: DATA COLLECTION

ELIGIBILITY IS THE PHENOMENON

Participants can recall a specific time AI output agreed with a judgment they had already formed. They do not need to have interpreted it in any particular way.

- Five or more years in audit or risk assurance, recurring engagements
- Criterion sampling plus referrals, all screened the same way; a short screening question tested first
- About 10 to 15, guided by the depth and richness of accounts
- Semi-structured, recorded interviews; protocol piloted first

6 · METHODOLOGY: DATA ANALYSIS

FIVE ACTIVITIES · NVIVO

- 1 Managing and organising. NVivo, named file plan, secure storage.
- 2 Reading and memoing. A dated reflexive journal, kept as evidence.
- 3 Describing and classifying. Horizontalisation, meaning units, clusters.
- 4 Developing interpretations. Textural and structural, then a composite.
- 5 Representing. A thematic map and a composite narrative.

7 · ETHICS AND VALIDATION

- Informed consent and IRB determination before any contact; pseudonyms, no firm detail
- Member checking of each participant's own textural description
- External reviewer reads purposefully chosen transcripts, including a disconfirming one, with the analytic memos
- Confirmation hazard log entries tied to specific analytic decisions in the audit trail

8 · DESIRED CONTRIBUTIONS

TO PRACTICE

Shows what human review of AI output is actually like, for firms and regulators relying on it.

TO ACADEMIA

Distinguishes the auditor's experience of agreement from measures of reliance.

9 · POTENTIAL LIMITATIONS

- Retrospective accounts of a past experience
- Insider researcher: bias managed, not eliminated
- Small purposive sample; not intended to generalise

AI DISCLOSURE

Generative AI (Claude, Anthropic) helped draft and lay out this poster from the author's proposal and protocol, and checked clarity. The research question, design decisions and interpretations are the author's (FIU Graduate School AI policy).

REFERENCES

Full list in the paper

Commerford, B. P., Dennis, S. A., Joe, J. R., & Ulla, J. W. (2022). Man versus machine: Complex estimates and auditor reliance on artificial intelligence. *Journal of Accounting Research*, 60(1), 171-201.

Kokina, J., Blanchette, S., Davenport, T. H., & Pachamanova, D. (2025). Challenges and opportunities for artificial intelligence in auditing. *International Journal of Accounting Information Systems*, 56, 100734.

Creswell, J. W., & Poth, C. N. (2024). *Qualitative inquiry and research design: Choosing among five approaches* (5th ed.). SAGE.

Murikah, W., Nthenge, J. K., & Musyoka, F. M. (2024). Bias and ethics of AI systems applied in auditing: A systematic review. *Scientific African*, 25, e02281.

Fotoh, L. E., & Mugwira, T. (2025). Exploring large language models in external audits: Implications and ethical considerations. *International Journal of Accounting Information Systems*, 56, 100748.

Parasuraman, R., & Manzey, D. H. (2010). Complacency and bias in human use of automation. *Human Factors*, 52(3), 381-410.

Glickman, M., & Sharot, T. (2025). How human-AI feedback loops alter human perceptual, emotional and social judgements. *Nature Human Behaviour*, 9(2), 345-359.

Tversky, A., & Kahneman, D. (1974). Judgment under uncertainty: Heuristics and biases. *Science*, 185(4157), 1124-1131.

Sampling and the interview

Every participant must have lived the phenomenon, not merely hold views about it.

ELIGIBILITY · ALL FOUR

1. Five or more years

in audit, internal audit or risk assurance

2. Recurring or multi-year engagements

where prior-period work is the anchoring context

3. Uses or has used AI or analytics tools

that produce conclusions, findings or risk ratings

4. Can recall a specific occasion

when such a tool reached a conclusion they had already reached

Criterion 4 is the study. A candidate who meets 1 to 3 and fails 4 has not lived the phenomenon, and their data would be opinion.

Maximum variation within the criterion:

Big Four, mid-tier and in-house internal audit; external and internal; senior associate to partner or CAE; several industries.

HOW MANY, AND WHEN TO STOP

Criterion sample of 10 to 15.

Guided by the depth and richness of the accounts, not saturation as an automatic rule (revised 1 Oct).

Screening question tested in the networks first; referrals screened the same way. Participants recall the occasion; they need not have interpreted it.

The stop decision, and its reasons, recorded in the audit trail on the day it is made.

Pilot first.

The protocol is piloted, and pilot testers cannot become participants.

Channel:

professional networks, association chapters and referral chains, not general research panels, which held about 6 eligible auditors per 100,000 in the author's prior study.

THE INTERVIEW · 60 MINUTES

Q1 (textural):

What have you experienced when an AI tool reached the same conclusion you had already reached?

Q2 (structural):

What situations or circumstances have typically influenced how you experienced that?

Grounding move:

before anything else, the participant silently recalls one specific occasion. If none comes, the screen failed and the interview ends courteously.

Body: the occasion (20 min) · prior-year conclusion versus AI answer (10) · confirmation as evidence (10) · when to stop looking (8) · under review (5).

Never asked:

"Why did you...?" (produces a defence) · "Would you say you were anchored?" (supplies the construct).

A NOTE FOR DR. GONZALEZ

Dr. Gonzalez, slide 1 is the poster I am presenting in Week 8, simplified as you suggested. These three appendix slides are for you only and are not presented; they set out the sampling, analysis and validation behind it, and slide 4 shows how I applied your 1 October feedback. The full detail is in the proposal paper. Thank you. Yasir Malik

Data analysis

Creswell and Poth's five analysis activities, carried out through Moustakas-style phenomenological steps in NVivo.

THE FIVE ACTIVITIES

1. Managing and organising the data.

Verbatim transcripts, pseudonyms at transcription, encrypted storage, one NVivo project.

2. Reading and memoing.

A reflexive memo before each listen-back, and emergent ideas memoed from the first transcript.

3. Describing and classifying codes into themes.

Horizontalisation, then meaning units, then clusters.

4. Developing and assessing interpretations.

A mandatory disconfirming pass before anything is written up.

5. Representing the data.

Textural, structural and composite descriptions, quoting rather than paraphrasing.

THE PHENOMENOLOGICAL STEPS

Epoché first.

Four written commitments, including the author's stake in the answer, set down before the first interview.

Horizontalisation.

Every significant statement listed, each given equal weight.

Meaning units and themes.

Statements reduced to non-overlapping units, then clustered.

Textural (what) and structural (how, in what context) descriptions,

for each participant and then across all of them.

Composite essence.

Two rules: no causal language, and no frequency claims such as "most participants".

THE AI BOUNDARY

AI may:

transcribe audio · format and tabulate what the author has already coded · check prose for clarity and grammar · act as an adversary ("what have I missed?") · find literature.

AI may not:

assign a code or meaning unit · name a theme · write any part of a description · decide what a participant meant · decide when to stop interviewing.

Why the line sits here:

the study examines what happens when a machine reaches a conclusion first and a human agrees with it. Letting a model propose themes and then agreeing with them would reproduce the phenomenon inside its own analysis.

Every AI use is logged with date, tool, request and outcome (FIU Graduate School AI policy; GEB 7911 limits).

Validation, and your feedback applied

Chapter 10: at least two strategies, drawn from more than one lens. This design uses seven.

THE RESEARCHER'S LENS

Reflexivity and the confirmation hazard log.

The author is a career auditor who built an AI review tool: insider access and insider bias at once. Every moment of feeling confirmed by a participant's account is logged, because that is the study's own phenomenon happening to the researcher.

Negative case analysis.

Actively seek auditors who distrusted or overrode the AI's agreement.

External review.

A few transcripts chosen on purpose, including one that challenges an emerging interpretation, read with the matching analytic memos.

THE PARTICIPANT'S AND THE READER'S LENS

Member checking.

Each participant receives their own textural description, not the transcript, to consider whether it represents what they experienced. Responses and how they were considered are recorded. It informs the individual descriptions, not the composite.

Thick description.

Firm type, function, engagement, tool category and the moment of agreement, so a reader can judge transfer.

Audit trail.

Protocol written before data collection, an append-only decision log, dated codebook versions and the hazard log: the study's own workpapers.

Falsification conditions

stated in advance of any participant.

YOUR FEEDBACK, APPLIED

1. Criterion kept.

Feasibility handled in recruitment: a tested screening question and referrals. Recall, not interpretation. Depth, not saturation.

2. Member checking scoped

to each individual textural description, with responses and their handling recorded.

3. External review purposeful:

a few chosen transcripts, one disconfirming, with memos, so the reviewer examines the move from accounts to meaning units, clusters and interpretations.

4. Hazard log works:

each entry records what I did about it, cross-referenced to the decision in the audit trail.

5. Poster simplified

and AI use disclosed on the slide.