

When the Machine Agrees: A Phenomenological Study of How Experienced Auditors Experience AI-Generated Confirmation of Judgments Already Formed

[NOTE: WORKING VERSION. Yellow text is a note to you; delete every yellow line before you submit. The title is a working title: change it if it does not sound like you.]

Yasir A. Malik

DBA Cohort 8.14

Florida International University, College of Business

GEB 7911 Qualitative Research Methods in Business, Dr. Cristina Gonzalez

October 6, 2026

Introduction

[NOTE: This section is YOUR Memo 4 text, carried over word for word. Each yellow flag sits under the sentence it applies to and says what Dr. Gonzalez (14 Sep) or Lakwantiera Braddy (13 Sep) asked for. Make the change in your own words. Target 3 to 4 double-spaced pages. The guidance asks this section to carry: a narrative hook, a short literature summary, the gap, the audiences, the purpose (script form), the RQ, the interpretive framework with four assumptions, and terminology.]

Focus and Research Problem

An auditor reviewing a multi-year engagement has always had a starting point: last year's conclusion. What has changed is that AI systems now introduce new risks into that process, including sycophancy and epistemic drift. These risks don't show up as one-time errors. They show up as a pattern that repeats every time the system is used.

[NOTE: Both praised this opening; keep it. Two small fixes: cite sycophancy here (Sharma et al., 2024) and either define 'epistemic drift' in Key Terminology or cut it. Write 'do not' rather than 'don't' in a formal paper.]

Professional standards assume a competent human reviewer sits above the AI system. Nobody has actually measured whether that reviewer's judgment survives contact with it. That matters because the auditors relying on those standards are the last check before a misjudgment reaches a client, a regulator, or a public filing.

[NOTE: SOFTEN (both reviewers, independently). 'Nobody has actually measured' is a categorical claim. Restate it as what the literature has mostly focused on. Cite the standard

you mean (PCAOB, 2024, on supervising generative AI output; NIST AI 600-1 calls this 'human-AI configuration').]

The empirical literature on AI and auditor judgment is thin, and what exists asks the wrong kind of question. Most work asks how much auditors rely on AI-generated output. That is a magnitude question, and a survey can answer it. It does not ask how auditors experience receiving that output, or how they make sense of it in the moment. That is a process question, and a survey cannot answer it. My own prior study tried to measure this by self-report and ran into the same limit every self-report measure of this kind runs into: it asks a person to notice the one thing the bias in question prevents them from noticing. The anchoring-and-adjustment literature originating with Tversky and Kahneman (1974) establishes that judgments formed from a starting point adjust insufficiently away from it, and that the effect persists even when participants are motivated toward accuracy. What that literature has not examined is what happens when the starting point is produced by a machine that also agrees with the conclusion the professional had already reached.

[NOTE: THREE FIXES. (1) 'asks the wrong kind of question' is the sentence both reviewers objected to. Suggested direction: the literature has focused primarily on the extent of reliance rather than the lived experience of receiving AI-generated confirmation.

(2) 'That is a process question' must go: process is grounded theory. Yours is a question about experience and meaning, which is phenomenology (Gonzalez 14 Sep; class 22 Sep).

(3) This paragraph is the literature summary the guidance asks for and it has one citation. Add 4 to 6 from the Literature Review list: Joyce and Biddle (1981) and Kinney and Uecker (1982) for anchoring in audit; Parasuraman and Manzey (2010) for automation

bias; Commerford et al. (2022) and Koreff (2022) for what has been measured about auditor reliance on AI.]

This matters to two audiences. For research, it separates two things the literature currently conflates: automation bias, which is a human tendency, and sycophancy, which is a property of the system. Asking how much someone relied on a machine cannot tell those two apart. For practice, audit firms and regulators are writing AI-use policies right now on the assumption that human review is a working control. Nobody has asked the people doing that reviewing what the experience of it actually is.

[NOTE: Gonzalez named the automation bias versus sycophancy distinction as a strength; keep it. SOFTEN 'Nobody has asked' (same issue as above). 'Conflates' also needs one citation, or soften it to 'often treats together'. Cite one policy document for 'writing AI-use policies' (PCAOB, 2024 or IAASB, 2024).]

Purpose Statement

The purpose of this phenomenological study is to understand the experience of receiving an AI-generated conclusion that confirms a judgment already formed, for experienced internal auditors, at organizations where AI-supported analytical tools are used in live engagement work. At this stage in the research, receiving AI-generated confirmation will be generally defined as the moment an AI system produces an analytical output matching a conclusion the auditor had already reached independently.

[NOTE: DECIDE ONE THING. This says 'internal auditors'. Your sampling plan and poster recruit across internal audit, external audit and risk assurance for maximum variation, and the RQ says 'experienced auditors'. Either drop 'internal' here

(recommended, it matches the RQ and the poster) or narrow the sampling. The Methodology draft assumes you drop it.]

Participants will be experienced internal auditors with eight or more years of practice, recurring multi-year engagement responsibility, and meaningful, direct exposure to AI-generated analytical output in the course of that work. Consistent with a phenomenological design, the site is defined by the shared professional context in which the phenomenon occurs rather than by a single named organization.

[NOTE: UPDATE. 'eight or more years' is now FIVE or more. Add the criterion Gonzalez asked for on 14 Sep: each participant must be able to recall a specific occasion when an AI tool reached a conclusion they had already reached. Recall is enough; they do not need to have interpreted it (her 1 Oct feedback).]

Research Question

How do experienced auditors experience and make sense of receiving an AI-generated conclusion that confirms a judgment they had already formed?

[NOTE: Braddy asked whether 'and make sense of' adds a second phenomenon. He did not insist. Add one sentence in your words saying why it stays: in phenomenology the meaning a person makes of an experience is part of the experience itself, not a separate process. Do not call it a process.]

Interpretive Framework and Assumptions

This study is approached from a social constructivist, inductive interpretive framework.

Ontological assumptions. Multiple realities exist. There is no single correct account of what it is like to receive machine-generated confirmation waiting to be recovered. There are as many accounts as there are auditors who have experienced it, and the variation among them is part of the finding rather than error to be averaged away.

Epistemological assumptions. Evidence consists of participants' own first-person descriptions of the experience, gathered through semi-structured, audio-recorded interviews. Closeness to participants is treated as a condition of knowing rather than a threat to it. The account is the data, not a proxy for something more objective sitting behind it.

Axiological assumptions. My own values are present in this study and are stated rather than bracketed away.

I am not an outsider to this setting. Fifteen years in audit and risk at Citigroup and JPMorgan Chase, and prior experience as a bank examiner for the Florida Office of Financial Regulation, give me direct familiarity with the vocabulary, pressures, and judgment calls this study asks participants to describe. That same background is also a risk: I already believe the phenomenon under study is real, which makes me more likely to hear confirmation in what a participant says than a genuinely neutral listener would. To guard against that, every occasion on which I notice myself feeling confirmed by a participant's account is logged in a confirmation hazard log kept for the duration of the study.

[NOTE: Both reviewers praised this paragraph; keep it. One change: remove 'Fifteen years' (your no-career-length rule). For example, start with 'My career in audit and risk at Citigroup and JPMorgan Chase...'. Add one clause that the hazard log records what you did about each entry, not only that you felt confirmed (Gonzalez, 1 Oct).]

Methodological assumptions. Of the five qualitative approaches, this study uses phenomenology. The research question asks what an experience is like and how it is made sense of, which is the question phenomenology is built to answer. A grounded theory design would commit the study to generating theory it is not yet positioned to generate, and a case study design would require a bounded site this study does not have.

Key Terminology

[NOTE: These six definitions came from Memo 4, where your disclosure says AI drafted them. Reword each in your own words and add one citation per term (anchoring: Tversky & Kahneman, 1974; automation bias: Parasuraman & Manzey, 2010; sycophancy: Sharma et al., 2024; professional judgment: Bonner, 2008). Add 'epistemic drift' if you keep it in the opening. Doing this also brings the paper's AI share down.]

Term	Definition
Anchoring	The tendency to rely too heavily on an initial reference point when forming a judgment, and to adjust insufficiently away from it.
Automation bias	The tendency to over-trust or defer to output from an automated system, even when independent evidence should prompt further scrutiny.
Sycophancy	A tendency in AI systems to agree with or affirm a user's stated position rather than challenge it. This is a property of the model, not a human bias.
Professional judgment	The application of relevant training, knowledge, and experience to reach a well-reasoned conclusion under conditions of uncertainty.
Recurring engagement	An audit relationship that continues across multiple annual periods with the same client, in which prior-period conclusions remain available as reference points.
AI-generated conclusion	An analytical output produced by an AI system, such as a flagged item, a risk score, or a suggested conclusion, that a reviewer can accept, adjust, or reject.

[NOTE: Memo 4 Section 6 (the larger programme) goes here, CONDENSED to two or three sentences (Braddy asked for this): the survey could not be fielded because the eligible population is about six per hundred thousand panel members; that same population can support 10 to 15 interviews. Replace the auditingai.github.io link, which is not live, with github.com/AuditingAI/Profile. DROP Memo 4 Section 7 (peer review questions); its answers now live in the Methodology below.]

Literature Review

[NOTE: YOU WRITE THIS SECTION. 3 to 4 double-spaced pages. Below is an outline and a source list only. Every source here already appears in the reference list of your research paper, so the citations are checked; you still need to read the ones you lean on before citing them. The example paper's pattern works well: for each theme, say what the literature shows, its strength, its weakness, and what that leaves open for this study. Close with a paragraph that states the gap in the softened form.]

2.1 Anchoring on a prior conclusion in audit judgment

- The question to answer: what do we already know about auditors carrying a starting point forward?
- Tversky & Kahneman (1974); Joyce & Biddle (1981); Kinney & Uecker (1982); Henrizi, Himmelsbach & Hunziker (2021); Furnham & Boo (2011)
- Mechanism worth naming: selective accessibility, where people test the anchor by looking for evidence that fits it (Strack & Mussweiler, 1997; Mussweiler & Strack, 1999). This is the closest thing in the literature to your 'confirmation' moment.
- Weakness to point out: almost all of it is experimental, and the anchor is a number or a prior-year figure, never a machine conclusion.

2.2 Automation bias: the human side

- Parasuraman & Manzey (2010) for complacency and automation bias; Dowling & Leech (2007) for audit decision aids.
- What it leaves open: the work studies whether people miss automation errors, not what it is like when the automation agrees with them and is right.

2.3 What has been measured about auditors and AI

- Commerford, Dennis, Joe & Ulla (2022): auditors discount AI evidence that CONTRADICTS management relative to the same evidence from a human specialist. Useful contrast: they studied disagreement; you study agreement.
- Koreff (2022): reliance on analytics conclusions varies with the input. Fedyk et al. (2022): firm-level AI investment and audit quality. Estep, Griffith & MacKenzie (2024): how financial executives respond to AI in audit.
- Field and practitioner evidence: Kokina et al. (2025); Fotoh & Mugwira (2025); Emett et al. (2025); Eulerich et al. (2024).
- CHECK BEFORE YOU WRITE THE GAP: open Kokina et al. (2025) and Fotoh & Mugwira (2025) and note their method. If either interviewed auditors about their experience, say so and narrow your gap to the specific moment of confirmation. This is the check your RB02 runbook asks for, and it is the first thing a reviewer will test.

2.4 Sycophancy: the system side

- Sharma et al. (2024) and Perez et al. (2023) on language models agreeing with the user; Fanous et al. (2025) on measuring it.
- Glickman & Sharot (2025) on human and AI feedback loops changing human judgment; Lee et al. (2025) on knowledge workers reporting less critical effort with generative AI; Messeri & Crockett (2024) on illusions of understanding.
- The point this subsection sets up: when the tool agrees, a human tendency and a system property point the same way, and a measure of reliance cannot tell them apart. That is your strongest original argument; make it in your words.

2.5 Professional skepticism and the reviewer

- Nelson (2009); Hurtt et al. (2013); Peecher, Solomon & Trotman (2013); Trotman & Yetton (1985) and Ramsay (1994) on review.
- Regulatory expectation that a human supervises AI output: PCAOB (2024); IAASB (2024); NIST (2024, AI 600-1).

2.6 Synthesis and the gap

- One paragraph. Suggested shape: the literature has examined the extent and direction of reliance, mostly by experiment and archive, and has examined sycophancy as a property of models; it has examined much less how experienced auditors experience the moment an AI tool confirms what they had already concluded. Keep it to what you have read.
- Optional: a small figure of the concepts the literature offers, with a line saying the analysis stays open to what participants describe (the example paper does this with its Figure 1).

Methodology

AI-assisted draft. Sections 4.3 to 4.7 were drafted with Claude (Anthropic) from my own research protocol and the instructor's feedback, then revised by me. Sections 4.1 and 4.2 are my own writing.

[NOTE: KEEP the italic label above in the final version: the course requires AI-written text to be labelled. Revise every paragraph in your voice; the more you rewrite, the lower the AI share. 4.1 and 4.2 are yours to write.]

4.1 Research Design

[NOTE: YOU WRITE THIS, half a page. It keeps the AI share under the 25% cap, and you have already written most of it: expand your 'Methodological assumptions' paragraph from the Introduction. Cover: phenomenology describes what people who lived the same experience share about it (Creswell & Poth, 2024; Moustakas, 1994); one line each on why not grounded theory (process and theory generation) and case study (bounded site); the site is the shared professional context, not one organization; transferability through rich description, no generalization.]

4.2 Role of the Researcher

[NOTE: YOU WRITE THIS. 3 to 4 paragraphs, first person. Most of it already exists in your own words: your axiology paragraph in the Introduction and the epoché in your protocol. Cover: (1) insider position, an asset and a risk, with no career-length number; (2) bracketing in practice: a reflexive memo before each listen-back, leading questions flagged in the transcript margin; (3) the confirmation hazard log and why feeling confirmed by a

participant is the phenomenon itself happening inside the study; (4) the insider and outsider pairing after Gioia and Chittipeddi (1991): their design, not their analysis method.

Name the risk that auditors may describe more careful practice to a researcher than they actually follow.]

4.3 Participants and Sampling

Participants will be selected by criterion sampling, because phenomenology requires that every participant has lived the experience under study rather than holding views about it (Creswell & Poth, 2024). To be eligible, a person must have five or more years in internal audit, external audit, or risk assurance; must have worked recurring engagements in which a prior-period conclusion is available as a reference point; must use or have used AI or automated analytical tools that produce conclusions, findings, or risk ratings; and must be able to recall a specific occasion when such a tool reached a conclusion they had already reached themselves. The last criterion carries the study. A candidate who meets the first three but cannot recall an occasion has not lived the phenomenon and is not enrolled.

Participants need to recall the experience, not to have interpreted it. They do not need to have thought of it as confirmation, reassurance, or bias. That meaning is for the interview to explore, so the screening question does not name or frame it. Before recruitment opens, the screening question will be tested informally with a small number of people in my professional network to learn whether experienced auditors can name a specific occasion. Those people will not become participants.

Recruitment will draw on my first-degree professional network and then on referrals. Every referred person passes the same screen. To keep variation across settings, the sample will

deliberately include different kinds of organizations, both internal and external audit, and a range of seniority and industries, with no more than two participants from any one firm. A commercial research panel will not be used. When the quantitative study was fielded, a panel of 334,976 members returned about twenty eligible people, which is too few for a survey but enough to recruit an interview study through a network.

The plan is 10 to 15 participants. The decision to stop will rest on the depth and richness of the accounts rather than on saturation as an automatic rule. As a guide, stopping will be considered when the accounts already gathered are full enough to support textural and structural descriptions, when three consecutive interviews add no new meaning unit, and when the external reviewer agrees. The judgment and its reasons will be recorded on the day it is made.

4.4 Data Collection

Data will come from one semi-structured interview per participant, of about 60 minutes, held by video call and audio-recorded with consent. The guide follows Creswell and Poth's (2024) two broad phenomenological questions: what the participant has experienced when an AI tool reached the conclusion they had already reached, and what situations or circumstances have typically shaped how they experienced it. Before the first question, each participant is asked to take a moment and bring one specific occasion to mind. The interview then follows that occasion: what they had concluded before the tool produced its output, how sure they were, what the tool said, what they did next, and what it was like to see it agree. Probes follow the participant rather than a fixed order.

Recordings will be transcribed verbatim. Pseudonyms are assigned at transcription and firm, client, and engagement names are removed. I will write a reflexive memo after each interview and before listening back to it.

4.5 Data Analysis

Analysis will follow the modified Stevick-Colaizzi-Keen method described by Moustakas (1994) and Creswell and Poth (2024), organized within the data analysis spiral of managing and organizing the data, reading and memoing, describing and classifying codes into themes, developing and assessing interpretations, and representing the data.

Managing and organizing. Transcripts and memos are held in NVivo on an encrypted device under a named file plan, one file per participant, with the audit trail kept alongside and nothing stored in any public location. **Reading and memoing.** Each transcript is read in full more than once, and memos record first impressions and any moment I notice myself agreeing with a participant. **Describing and classifying.** Every statement about the experience is listed with equal weight (horizontalization), reduced to non-repeating meaning units in language close to the participant's own, and grouped into themes named in words a participant would recognize. A unit that does not fit is set aside in a visible list rather than dropped. **Developing and assessing interpretations.** A textural description sets out what participants experienced, quoted throughout, and a structural description sets out the conditions in which they experienced it, such as deadline pressure, the review hierarchy, and engagement history. Before the composite description of the essence is written, every transcript is re-read only for material that contradicts it, including accounts of checking harder when the tool agreed. **Representing.** Findings will be written as a narrative with extended quotation, without frequency counts or causal claims.

Coding will be done by me. NVivo stores and organizes the codes I assign; its automated coding features will not be used, and AI tools will not assign codes, name themes, write descriptions, or decide what a participant meant. Letting a model propose themes and then agreeing with them would reproduce, inside the analysis, the experience the study examines. Any use of AI, for example checking prose for clarity, is logged with the date, the tool, and what was done with the output.

4.6 Trustworthiness and Validation

Trustworthiness is addressed through Lincoln and Guba's (1985) criteria and Creswell and Miller's (2000) validation strategies, using the participant, reviewer, and researcher lenses.

Member checking. Each participant receives their own textural description, not the transcript, as a chance to consider whether it represents what they experienced and to add context. Responses and how each was considered are recorded. Member checking informs the credibility of the individual descriptions; it does not validate the structural interpretations or the composite description. **External review.** A reviewer with no stake in the argument will examine a few transcripts chosen on purpose, including one that challenges an emerging interpretation, together with the matching analytic memos, so that the review looks at how I move from accounts to meaning units, themes, and interpretations. **Researcher reflexivity.** The confirmation hazard log records each occasion on which I felt confirmed by an account and what I did about it, for example returning to the transcript for contradicting material or rewording a probe for the next interview. Each entry is cross-referenced in the audit trail to the decision it affected. **Audit trail.** The protocol was written and dated before any data collection, and each design decision is logged as it is made.

4.7 Ethical Considerations

No participant will be approached before IRB approval. The existing approval (IRB-25-0462) covers an anonymous online survey, so an interview study with audio recording requires a modification or a new protocol, and the IRB will be asked which applies. Consent will be recorded verbally, with separate consent for audio recording. The main risk to participants is professional: some may describe practices their organization would not endorse. Pseudonyms, removal of identifying detail at transcription, encrypted storage, deletion of recordings after transcription and member checking, and the right to withdraw any passage without giving a reason all address that risk. No recording, transcript, or identifier will be uploaded to a public AI platform, consistent with the FIU University Graduate School AI policy.

Expected Contributions

[NOTE: YOU WRITE THIS. One page. The points below are prompts, not text. Pick three or four and argue them in your own voice; this is where the course most wants to hear you.]

- Scholarship: describes the lived experience of AI confirmation, which experimental and archival work on reliance does not reach.
- Scholarship: separates automation bias (human) from sycophancy (system) in the one situation where they point the same way and cannot be told apart by measuring reliance.
- Practice: tests from the inside the assumption in firm and regulatory AI policies that human review is a working control; what reviewers say they do when the tool agrees could inform review guidance and training.
- Method: the confirmation hazard log as a reflexive control for any study where the researcher holds a prior position, shown changing decisions rather than only recording them.
- Programme: this is the second link of your three-link chain. Say in one line how its findings will feed the dissertation's quantitative work.
- Limitations to name briefly: a small, network-recruited sample; self-report of a moment that may have been tidied in memory; possible description of more careful practice than is followed; no generalization claimed.

Proposed Timeline

AI-assisted draft, built from my research protocol and revised by me.

The study is planned to run from October 2026 to June 2027. Data collection and analysis run in parallel, so each interview is memoed and coded before the next is held.

Phase	Activities	Period
1. Approval	IRB modification or new protocol; consent and recording language; confirm transcription service and external reviewer	October to December 2026
2. Preparation	Test the screening question informally; finalize the interview guide; recruit the external reviewer	December 2026
3. Recruitment	Network outreach and referrals; screening; scheduling	January to February 2027
4. Interviews and early analysis	10 to 15 interviews; transcription and reflexive memos; horizontalization and meaning units as each transcript is ready; hazard log kept	January to March 2027
5. Analysis	Themes; textural and structural descriptions; disconfirming re-read; decision on stopping recorded	March to April 2027
6. Validation	Member checking on individual descriptions; external review of selected transcripts with memos	April to May 2027
7. Write-up	Composite description; findings narrative; limitations; audit trail closed	May to June 2027

[NOTE: Check the dates against your own calendar and the dissertation work. Do not add a completion date for the DBA anywhere in this paper.]

References

[NOTE: Every entry below is from your research paper's verified reference list or a standard methods text. Delete any you do not end up citing, and add page numbers for direct quotations.]

Bonner, S. E. (2008). *Judgment and decision making in accounting*. Prentice Hall.

Commerford, B. P., Dennis, S. A., Joe, J. R., & Ulla, J. W. (2022). Man versus machine: Complex estimates and auditor reliance on artificial intelligence. *Journal of Accounting Research*, 60(1), 171-201. <https://doi.org/10.1111/1475-679X.12407>

Creswell, J. W., & Miller, D. L. (2000). Determining validity in qualitative inquiry. *Theory Into Practice*, 39(3), 124-130.

Creswell, J. W., & Poth, C. N. (2024). *Qualitative inquiry and research design: Choosing among five approaches* (5th ed.). SAGE Publications.

Dowling, C., & Leech, S. A. (2007). Audit support systems and decision aids. *International Journal of Accounting Information Systems*, 8(2), 92-116.

Emett, S. A., Eulerich, M., Lipinski, E., Prien, N., & Wood, D. A. (2025). Leveraging ChatGPT for enhancing the internal audit process: A real-world example from Uniper, a large multinational company. *Accounting Horizons*, 39(2), 125-135. <https://doi.org/10.2308/HORIZONS-2023-111>

Estep, C., Griffith, E. E., & MacKenzie, N. L. (2024). How do financial executives respond to the use of artificial intelligence in financial reporting and auditing? *Review of Accounting Studies*, 29(3), 2798-2831. <https://doi.org/10.1007/s11142-023-09771-y>

- Eulerich, M., Sanatizadeh, A., Vakilzadeh, H., & Wood, D. A. (2024). Is it all hype? ChatGPT's performance and disruptive potential in the accounting and auditing industries. *Review of Accounting Studies*, 29(3), 2318-2349. <https://doi.org/10.1007/s11142-024-09833-9>
- Fanous, A., Goldberg, J., Agarwal, A. A., Lin, J., Zhou, A., Daneshjou, R., & Koyejo, S. (2025). *SycEval: Evaluating LLM sycophancy* (arXiv:2502.08177). arXiv. <https://doi.org/10.48550/arXiv.2502.08177>
- Fedyk, A., Hodson, J., Khimich, N., & Fedyk, T. (2022). Is artificial intelligence improving the audit process? *Review of Accounting Studies*, 27(3), 938-985. <https://doi.org/10.1007/s11142-022-09697-x>
- Fotoh, L. E., & Mugwira, T. (2025). Exploring large language models in external audits: Implications and ethical considerations. *International Journal of Accounting Information Systems*, 56, 100748. <https://doi.org/10.1016/j.accinf.2025.100748>
- Furnham, A., & Boo, H. C. (2011). A literature review of the anchoring effect. *The Journal of Socio-Economics*, 40(1), 35-42.
- Gioia, D. A., & Chittipeddi, K. (1991). Sensemaking and sensegiving in strategic change initiation. *Strategic Management Journal*, 12(6), 433-448.
- Glickman, M., & Sharot, T. (2025). How human-AI feedback loops alter human perceptual, emotional and social judgements. *Nature Human Behaviour*, 9(2), 345-359. <https://doi.org/10.1038/s41562-024-02077-2>

- Henrizi, P., Himmelsbach, D., & Hunziker, S. (2021). Anchoring and adjustment effects on audit judgments: Experimental evidence from Switzerland. *Journal of Applied Accounting Research*, 22(4), 598-621. <https://doi.org/10.1108/JAAR-01-2020-0011>
- Hurt, R. K., Brown-Liburd, H., Earley, C. E., & Krishnamoorthy, G. (2013). Research on auditor professional skepticism: Literature synthesis and opportunities for future research. *Auditing: A Journal of Practice & Theory*, 32(Supplement 1), 45-97.
- International Auditing and Assurance Standards Board. (2024). *Technology position and focus area resources*. <https://www.iaasb.org/focus-areas/technology>
- Joyce, E. J., & Biddle, G. C. (1981). Anchoring and adjustment in probabilistic inference in auditing. *Journal of Accounting Research*, 19(1), 120-145.
- Kinney, W. R., Jr., & Uecker, W. C. (1982). Mitigating the consequences of anchoring in auditor judgments. *The Accounting Review*, 57(1), 55-69.
- Kokina, J., Blanchette, S., Davenport, T. H., & Pachamanova, D. (2025). Challenges and opportunities for artificial intelligence in auditing: Evidence from the field. *International Journal of Accounting Information Systems*, 56, 100734. <https://doi.org/10.1016/j.accinf.2025.100734>
- Koreff, J. (2022). Are auditors' reliance on conclusions from data analytics impacted by different data analytic inputs? *Journal of Information Systems*, 36(1), 19-37. <https://doi.org/10.2308/ISYS-19-051>
- Lee, H.-P., Sarkar, A., Tankelevitch, L., Drosos, I., Rintel, S., Banks, R., & Wilson, N. (2025). The impact of generative AI on critical thinking: Self-reported reductions in cognitive

effort and confidence effects from a survey of knowledge workers. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*.
<https://doi.org/10.1145/3706598.3713778>

Lincoln, Y. S., & Guba, E. G. (1985). *Naturalistic inquiry*. SAGE Publications.

Messeri, L., & Crockett, M. J. (2024). Artificial intelligence and illusions of understanding in scientific research. *Nature*, 627(8002), 49-58. <https://doi.org/10.1038/s41586-024-07146-0>

Moustakas, C. (1994). *Phenomenological research methods*. SAGE Publications.

Mussweiler, T., & Strack, F. (1999). Hypothesis-consistent testing and semantic priming in the anchoring paradigm: A selective accessibility model. *Journal of Experimental Social Psychology*, 35(2), 136-164.

National Institute of Standards and Technology. (2024). *Artificial intelligence risk management framework: Generative artificial intelligence profile* (NIST AI 600-1). U.S. Department of Commerce. <https://doi.org/10.6028/NIST.AI.600-1>

Nelson, M. W. (2009). A model and literature review of professional skepticism in auditing. *Auditing: A Journal of Practice & Theory*, 28(2), 1-34.

Parasuraman, R., & Manzey, D. H. (2010). Complacency and bias in human use of automation: An attentional integration. *Human Factors*, 52(3), 381-410.
<https://doi.org/10.1177/0018720810376055>

- Peecher, M. E., Solomon, I., & Trotman, K. T. (2013). An accountability framework for financial statement auditors and related research questions. *Accounting, Organizations and Society*, 38(8), 596-620.
- Perez, E., Ringer, S., Lukošīūtė, K., Nguyen, K., Chen, E., Heiner, S., et al. (2023). Discovering language model behaviors with model-written evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023* (pp. 13387-13434). <https://doi.org/10.18653/v1/2023.findings-acl.847>
- Public Company Accounting Oversight Board. (2024). *Spotlight: Staff update on outreach activities related to the integration of generative artificial intelligence in audits and financial reporting*. <https://pcaobus.org/documents/generative-ai-spotlight.pdf>
- Ramsay, R. J. (1994). Senior/manager differences in audit workpaper review performance. *Journal of Accounting Research*, 32(1), 127-135.
- Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Asbell, A., Bowman, S. R., et al. (2024). Towards understanding sycophancy in language models. In *Proceedings of the Twelfth International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/2310.13548>
- Strack, F., & Mussweiler, T. (1997). Explaining the enigmatic anchoring effect: Mechanisms of selective accessibility. *Journal of Personality and Social Psychology*, 73(3), 437-446.
- Trotman, K. T., & Yetton, P. W. (1985). The effect of the review process on auditor judgments. *Journal of Accounting Research*, 23(1), 256-267.

Tversky, A., & Kahneman, D. (1974). Judgment under uncertainty: Heuristics and biases.
Science, 185(4157), 1124-1131.

Disclosure of Artificial Intelligence Use

[NOTE: Required by the FIU University Graduate School AI policy and the course's labelling rule. Update it to match what you actually end up keeping.]

Generative AI (Claude, Anthropic) was used in preparing this proposal in four ways. It drafted the Participants and Sampling, Data Collection, Data Analysis, Trustworthiness, and Ethical Considerations sections of the Methodology, and the Proposed Timeline, from decisions already recorded in my own research protocol and from the instructor's feedback; those sections are labelled in the text and I revised them. It organized an outline and a source list for the Literature Review and Expected Contributions, which I wrote. It compiled the reference list from sources I had already gathered, and it formatted the document. The research problem, purpose statement, research question, interpretive framework and assumptions, literature review, research design, role of the researcher, expected contributions, and every methodological decision are my own. No participant data exists, and none was shared with any AI tool.