



Receiving Confirmation

How experienced auditors make sense of AI-generated agreement with a judgment they had already formed

CENTRAL RESEARCH QUESTION

How do experienced auditors experience and make sense of receiving an AI-generated conclusion that confirms a judgment they had already formed?

1 · FOCUS AND RESEARCH PROBLEM

An auditor on a recurring engagement has always carried a prior conclusion forward. What is new is that an AI system now produces a conclusion of its own, and sometimes it agrees.

Research on AI and auditor judgment has focused primarily on the extent of reliance, which a survey can measure. What remains unexplored is the lived experience of receiving machine agreement with a judgment already formed.

Two things are being conflated: automation bias, a human tendency, and sycophancy, a property of the system. A magnitude-of-reliance measure cannot tell them apart.

Standards assume a competent human reviewer sits above the system. Firms and regulators are writing AI-use policy on that assumption now.

2 · RESEARCH PURPOSE

The purpose of this phenomenological study is to understand how experienced auditors experience and make sense of receiving an AI-generated conclusion that confirms a judgment they had already formed. The phenomenon under study is the experience of encountering machine-generated confirmation of a prior professional judgment.

3 · FRAMEWORK AND ASSUMPTIONS

SOCIAL CONSTRUCTIVIST · INDUCTIVE

ONTOLOGICAL

Multiple realities. Auditors' accounts of the experience are the data, not a proxy for one correct account.

EPISTEMOLOGICAL

First-person accounts, gathered in semi-structured audio-recorded interviews and read in the participant's own words before coding.

AXIOLOGICAL

A career in bank examination, regulatory reporting and internal audit, including building an AI review tool. Insider access and insider bias at once. Every occasion of feeling confirmed is entered in a confirmation hazard log.

METHODOLOGICAL

Phenomenology. A phenomenon shared across individuals at different firms, with no single site to enter.

4 · QUALITATIVE APPROACH

Phenomenology (Creswell & Poth, 2024). The question asks about experience and meaning making across individuals, which is what phenomenology is for.

- Case study needs a bounded site; this phenomenon crosses firms
- Grounded theory studies process; this study asks about experience
- Ethnography needs embedded access to a firm's review process

5 · DATA COLLECTION

ELIGIBILITY IS THE PHENOMENON

A participant must have received AI-generated output that confirmed a judgment they had already formed independently. General experience of AI-supported audit work does not qualify.

- Five or more years in audit or risk assurance, with recurring engagement responsibility
- Criterion sampling, 10 to 15. Saturation is a stopping rule, not a target
- No single site. Professional networks, association chapters, referral chains
- Semi-structured interviews, recorded and transcribed verbatim. Protocol piloted first, and pilot testers cannot become participants
- Encrypted storage, pseudonyms at transcription, named file plan

6 · DATA ANALYSIS

FIVE ACTIVITIES · NVIVO

- 1 Managing and organising. NVivo, named file plan, secure storage.
- 2 Reading and memoing. A dated reflexive journal, kept as evidence.
- 3 Describing and classifying. Horizontalisation, meaning units, clusters.
- 4 Developing interpretations. Textural and structural, then a composite.
- 5 Representing. A thematic map and a composite narrative.

7 · ETHICAL CONSIDERATIONS

- Written informed consent before recording begins
- IRB determination confirmed before any contact, not assumed
- Pseudonyms at transcription. No client, engagement or firm-identifying detail
- Withdrawal at any point, data destroyed on request
- Member checking of the composite description offered to every participant
- No confidential data to any public AI platform. AI use disclosed per FIU policy

8 · EXPECTED CONTRIBUTION

TO PRACTICE

Firms and regulators are writing AI-use policy on the assumption that human review is a working control. This describes what that review is actually like from inside it.

TO ACADEMIA

It separates automation bias from sycophancy, which measures of how much someone relied on a machine cannot distinguish.

9 · POTENTIAL LIMITATIONS

- Retrospective self-report. Participants describe an experience after the fact
- Insider researcher. Managed by the confirmation hazard log, not eliminated
- Context-dependent. Not intended to generalise
- One national standards environment, small purposive sample

TIMELINE

M1

IRB determination, protocol pilot

M2

Recruitment and consent

M3 to M5

Interviews, transcription, memoing

M5 to M7

Coding, clustering, composite

M8

Member checking and write-up

REFERENCES

Full list in the paper

Commerford, B. P., Dennis, S. A., Joe, J. R., & Ulla, J. W. (2022). Man versus machine: Complex estimates and auditor reliance on artificial intelligence. *Journal of Accounting Research*, 60(1), 171-201.

Kokina, J., Blanchette, S., Davenport, T. H., & Pachamanova, D. (2025). Challenges and opportunities for artificial intelligence in auditing. *International Journal of Accounting Information Systems*, 56, 100734.

Creswell, J. W., & Poth, C. N. (2024). *Qualitative inquiry and research design: Choosing among five approaches* (5th ed.). SAGE.

Murikah, W., Nthenge, J. K., & Musyoka, F. M. (2024). Bias and ethics of AI systems applied in auditing: A systematic review.

Fotoh, L. E., & Mugwira, T. (2025). Exploring the impact of generative AI on professional skepticism in auditing.

Parasuraman, R., & Manzey, D. H. (2010). Complacency and bias in human use of automation. *Human Factors*, 52(3), 381-410.

Glickman, M., & Sharot, T. (2025). How human-AI feedback loops alter human perceptual, emotional and social judgements.

Tversky, A., & Kahneman, D. (1974). Judgment under uncertainty: Heuristics and biases. *Science*, 185(4157), 1124-1131.

Sampling and the interview

Every participant must have lived the phenomenon, not merely hold views about it.

ELIGIBILITY · ALL FOUR

1. Five or more years

in audit, internal audit or risk assurance

2. Recurring or multi-year engagements

where prior-period work is the anchoring context

3. Uses or has used AI or analytics tools

that produce conclusions, findings or risk ratings

4. Can recall a specific occasion

when such a tool reached a conclusion they had already reached

Criterion 4 is the study. A candidate who meets 1 to 3 and fails 4 has not lived the phenomenon, and their data would be opinion.

Maximum variation within the criterion:

Big Four, mid-tier and in-house internal audit; external and internal; senior associate to partner or CAE; several industries.

HOW MANY, AND WHEN TO STOP

Criterion sample of 10 to 15.

Saturation is a stopping rule, not a target.

Stop after three consecutive interviews that add no new surviving meaning unit.

Below 10, do not stop regardless of apparent saturation. Above 15, stop and report why saturation did not arrive, which is itself a finding.

Pilot first.

The protocol is piloted, and pilot testers cannot become participants.

Channel:

professional networks, association chapters and referral chains, not general research panels, which held about 6 eligible auditors per 100,000 in the author's prior study.

THE INTERVIEW · 60 MINUTES

Q1 (textural):

What have you experienced when an AI tool reached the same conclusion you had already reached?

Q2 (structural):

What situations or circumstances have typically influenced how you experienced that?

Grounding move:

before anything else, the participant silently recalls one specific occasion. If none comes, the screen failed and the interview ends courteously.

Body: the occasion (20 min) · prior-year conclusion versus AI answer (10) · confirmation as evidence (10) · when to stop looking (8) · under review (5).

Never asked:

"Why did you...?" (produces a defence) · "Would you say you were anchored?" (supplies the construct).

A NOTE FOR DR. GONZALEZ

Dr. Gonzalez, this is a working draft, shared now so your feedback can shape the final version. Slide 1 is the one-page poster I will present in Week 8. These three appendix slides are not for presentation; they set out the sampling, analysis and validation behind it, and slide 4 ends with four questions where your guidance would help most. I will send the revised poster and proposal next week. Thank you.
Yasir Malik

Data analysis

Creswell and Poth's five analysis activities, carried out through Moustakas-style phenomenological steps in NVivo.

THE FIVE ACTIVITIES

1. Managing and organising the data.

Verbatim transcripts, pseudonyms at transcription, encrypted storage, one NVivo project.

2. Reading and memoing.

A reflexive memo before each listen-back, and emergent ideas memoed from the first transcript.

3. Describing and classifying codes into themes.

Horizontalisation, then meaning units, then clusters.

4. Developing and assessing interpretations.

A mandatory disconfirming pass before anything is written up.

5. Representing the data.

Textural, structural and composite descriptions, quoting rather than paraphrasing.

THE PHENOMENOLOGICAL STEPS

Epoché first.

Four written commitments, including the author's stake in the answer, set down before the first interview.

Horizontalisation.

Every significant statement listed, each given equal weight.

Meaning units and themes.

Statements reduced to non-overlapping units, then clustered.

Textural (what) and structural (how, in what context) descriptions,

for each participant and then across all of them.

Composite essence.

Two rules: no causal language, and no frequency claims such as "most participants".

THE AI BOUNDARY

AI may:

transcribe audio · format and tabulate what the author has already coded · check prose for clarity and grammar · act as an adversary ("what have I missed?") · find literature.

AI may not:

assign a code or meaning unit · name a theme · write any part of a description · decide what a participant meant · decide when saturation is reached.

Why the line sits here:

the study examines what happens when a machine reaches a conclusion first and a human agrees with it. Letting a model propose themes and then agreeing with them would reproduce the phenomenon inside its own analysis.

Every AI use is logged with date, tool, request and outcome (FIU Graduate School AI policy; GEB 7911 limits).

Validation, and where feedback would help most

Chapter 10: at least two strategies, drawn from more than one lens. This design uses seven.

THE RESEARCHER'S LENS

Reflexivity and the confirmation hazard log.

The author is a career auditor who built an AI review tool: insider access and insider bias at once. Every moment of feeling confirmed by a participant's account is logged, because that is the study's own phenomenon happening to the researcher.

Negative case analysis.

Actively seek auditors who distrusted or overrode the AI's agreement.

Outsider review.

Every third transcript read cold by a second person asking "where is the researcher in this data?"

THE PARTICIPANT'S AND THE READER'S LENS

Member checking.

Each participant receives their own textural description, not the raw transcript, and is asked what is wrong or missing.

Thick description.

Firm type, function, engagement, tool category and the moment of agreement, so a reader can judge transfer.

Audit trail.

Protocol written before data collection, an append-only decision log, dated codebook versions and the hazard log: the study's own workpapers.

Falsification conditions

stated in advance of any participant.

QUESTIONS FOR DR. GONZALEZ

1. Is criterion 4 too narrow?

Requiring a recalled occasion of AI agreement may make 10 participants hard to reach. Would you keep it strict, or allow a closely related occasion?

2. Member checking on the textural description

rather than the transcript: acceptable for this proposal, or should the transcript also go back to participants?

3. One outsider reviewer

reading every third transcript: enough, given how close the author is to the phenomenon?

4. Is the AI boundary drawn in the right place

for the course's 25% limit and labelling rule?