Source 2 — The dissertation argument, and every challenge to it

*The AI extension. This is the material to be able to **defend**, not merely recite.

Chain version 1.1, current as of 11 August 2026.*

---

1. The chain

```
AI tool enters recurring review work
↓ L1
system output becomes the anchor — automated, continuous,
arriving before the reviewer has formed a view
↓ L2
the model confirms the position the reviewer has already stated
rather than challenging it
↓ L3
successive models reprocess the same work; conclusions converge
on each other rather than on evidence
↓
recurring professional judgment degrades
while every individual model passes its own validation
```

That last line is the thesis. **The failure is invisible to model validation** because no single model is malfunctioning. Each performs to specification. The degradation lives in the sequence between the model and the human, and in the loop between models — neither of which any validation plan currently inspects.

2. Why each link is claimed, and where each is weak

L1 — system output becomes the anchor

The claim. When an AI tool produces a conclusion before the reviewer forms their own, that output functions as an anchor. Unlike a prior-year figure, it is generated automatically, continuously, and at scale.

Support. Anchoring-and-adjustment is among the most replicated findings in judgment research, and the qualifying study's entire model rests on it.

The strongest challenge. Classic anchoring effects have taken replication damage in adjacent literatures, and effect sizes in expert populations run smaller than in the undergraduate samples that produced the canonical results.

Status: defensible — but only if argued with expert-population evidence. Citing undergraduate lab studies to a committee of accounting academics is the weakest possible version of this argument.

L2 — the model confirms rather than challenges

The claim. Language models tend to agree with a position the user has already stated instead of contesting it. A reviewer who says "this looks reasonable" gets agreement, not challenge.

Support. Sycophancy is documented and named in the LLM literature; vendors publish on mitigating it, which is itself evidence it exists.

The strongest challenge. Sycophancy is a moving target. Each model generation is explicitly trained against it. **A claim true in 2026 may be false at a 2028 defense.**

Status: time-sensitive, and must always carry a date. The defensible framing is *mechanism, not artifact*: sycophancy arises from optimising on human approval, and that optimisation pressure persists across generations even as each generation is patched. Argue the pressure, not the symptom.

L3 — models converge on each other

The claim. As successive models reprocess work that earlier models produced, conclusions converge on one another rather than on the underlying evidence.

Support. Model-collapse and synthetic-data-contamination results. **And, since August 2026, external support:** *How LLMs Audit Each Other: Five Mechanisms of Auditor Bias in Cross-Model Peer Review Under Identity Disclosure and Cross-Lingual Conditions* (2026). Cross-model peer review **is** this link, instrumented by someone else, with five named mechanisms.

The strongest challenge. Nearly all of that evidence concerns *training corpora* — models degrading when trained on model output. The step from "models degrade on synthetic data" to "auditors converge because the models they consult converge" is **an inference, not a finding.**

Status: weakest link. This is where a committee pushes. It is also the largest opportunity — see §4.

The whole chain

Status: untested. Say so. No data exists. Nothing in this argument is a finding.

3. Three standing challenges — carry these until answered

1 · The improvement case. A 2026 study found AI assistance *reduced* bias in a judgment task (news summarisation, lay readers). Different domain, non-professional task — so not fatal. But an argument that AI degrades professional judgment has to engage the case that it sometimes improves it. If it doesn't, a committee member will raise it first, and the argument will look one-sided.

2 · The moving-target problem. See L2. Any claim about model behaviour needs a date attached and a mechanism underneath it.

3 · L3's inferential gap. Model-to-model convergence is documented. Auditor-to-auditor convergence *via* models is not.

4. Where the contribution actually is

The weakest link is also the most valuable one, because it is the one a well-designed study could **measure**.

Nobody has shown that professionals reading model output converge on each other. That is an empirical question with a tractable design: give reviewers the same case, vary whether they see model output and whose output they see, and measure dispersion of conclusions.

The dissertation's contribution is closing L3's gap, not restating the chain. Restating it is commentary. Measuring it is research.

5. Practice questions — try these before opening the sources

1. A committee member says: *"Anchoring has replication problems. Why should I believe your L1?"*
2. *"You're describing 2026 model behaviour. Your defense is in 2028. Why won't this be obsolete?"*
3. *"Model collapse is about training data. Your claim is about people. Where's the bridge?"*
4. *"There's evidence AI reduces bias. Why does your chain only run one direction?"*
5. *"If every model passes validation, why should a regulator care?"*
6. *"Which single link, if false, collapses the whole argument?"*
7. *"What would falsify your chain?"*

Question 7 is the one to be able to answer instantly. An argument with no falsification condition is not a research argument.

6. Change log

- **v1.1 — 10 Aug 2026.** L3 upgraded from conjecture to weakest-link-with-external-support after the cross-model peer-review paper. Improvement-case challenge added. L2 reframed as mechanism.
- **v1.0 — Jul 2026.** Chain first stated in the qualifying manuscript. Praised as sound research judgment; no link individually defended.

The chain is re-tested at every literature sweep. A link that survives unchanged still gets its date bumped, so "defensible" never quietly means "unexamined since last year."